

NN-Adaptive Bayesian Quantum Estimation

Michał Chrzanowski

7 September 2026

Learning **when to measure** a qubit, so that you learn **what its frequency is** as fast as physically possible.

This repository implements adaptive Bayesian experimental design for single-qubit Ramsey frequency estimation. A sequential Monte Carlo filter maintains a posterior over the unknown Larmor frequency ω ; a policy reads that posterior and picks the interrogation time t for the next shot; the measurement outcome sharpens the posterior; repeat. The policy is trained by reinforcement learning (TRPO) or by a gradient-free evolutionary search (CEM), following Fiderer, Schuff & Braun, *Neural-network heuristics for adaptive Bayesian quantum estimation* — the MLflow runs are tagged `project: fiderer`.

On the reference. Equation numbers throughout this README refer to `fiderer.pdf` in this repository, which is the preprint (dated 2020-03-05); numbering may differ from the published version, PRX Quantum **2**, 020303 (2021). The physics below has been checked against that PDF equation by equation. Where this implementation departs from the paper — and it does, in ways that matter — it is catalogued under Divergences from Fiderer et al.

1 Table of contents

- The physics
- The inference layer
- The control problem
- Methods implemented
- Repository layout
- Installation
- Running things
 - TRPO, end to end
 - CEM, end to end
 - Baselines (no training)
- Experiment tracking

- Baseline results
- Validation over 10 000 omega
- Divergences from Fiderer et al.
- Known issues and gotchas

2 The physics

2.1 The experiment

A single qubit undergoing **Ramsey interferometry** in the presence of pure dephasing:

1. Prepare $|+\rangle = (|0\rangle + |1\rangle)/\sqrt{2}$.
2. Let it precess freely under $H = \frac{\omega}{2}\sigma_z$ for a time t . This is the **control variable** — the one thing the experimenter chooses.
3. Apply a $\pi/2$ read-out pulse and measure in the computational basis.
4. Record the binary outcome $d \in \{0, 1\}$.

The qubit is not isolated: transverse noise destroys the coherence $\rho_{01}(t) = \rho_{01}(0) e^{-t/T_2} e^{-i\omega t}$ on a timescale T_2 .

2.2 The likelihood

Solving the Lindblad pure-dephasing master equation and applying the Born rule gives the outcome probability implemented in `modules/simulation.py:13-16` (NumPy) and `modules/simulation.py:51-53` (Torch):

$$P(d = 0 \mid \omega, t) = e^{-t/T_2} \cos^2\left(\frac{\omega t}{2}\right) + \frac{1 - e^{-t/T_2}}{2} \equiv \frac{1}{2} \left[1 + e^{-t/T_2} \cos(\omega t) \right]$$

and $P(d = 1) = 1 - P(d = 0)$. The two forms are algebraically identical.

Read it as a **decaying interference fringe**:

term	meaning
$\cos^2(\omega t/2)$	the Ramsey fringe — this is what carries information about ω
$v(t) = e^{-t/T_2}$	the visibility , or coherent fraction
$(1 - v)/2$	the fully dephased fraction, which returns a coin flip and tells you nothing

As $t \rightarrow \infty$ the visibility vanishes, $P(d = 0) \rightarrow 1/2$, and the measurement becomes pure noise. **Decoherence therefore imposes a hard ceiling on how long it is worth interrogating.**

Both the simulator and the filter's likelihood use this same expression, and T_2 is treated as exactly known. The model is well-specified: there is no model mismatch and no nuisance parameter.

2.3 What is being estimated

quantity	value	where
parameter	ω , a single scalar (the Larmor frequency / detuning)	<code>seq_montecarlo.py:9</code>
prior	$\omega \sim \text{Uniform}(0, 1)$, so $\text{Var}[\omega] = 1/12 \approx 0.0833$	<code>seq_montecarlo.py:9, :127</code>
true values	drawn from the same $\text{Uniform}(0, 1)$	<code>train_TRPO_baseline.py:67</code>
coherence time	$T_2 = 10$	<code>modules/simulation.py:31</code> (see gotchas)
control	$t \in [0.1, 3000]$	<code>train_TRPO_baseline.py:54-55</code>

Units are arbitrary — ω and t only ever appear as the product ωt , and T_2 shares the units of t . With $\omega \lesssim 1$ and $T_2 = 10$, roughly $\omega T_2 \lesssim 10$ radians of phase accumulate before coherence is lost, so a few fringes are resolvable within the coherence window.

2.4 The figure of merit: Bayes risk

This is a **Bayesian** estimation problem, and the quantity everything is judged by is the **Bayes risk** — Fiderer Eq. (1):

$$r(h \mid p(\theta)) = \mathbb{E}_{D_k} \left[\text{tr Cov}(\theta \mid D_k) \right]$$

where h is the experiment-design heuristic, $D_k = (d_1, \dots, d_k)$ the measurement record, and smaller is better. It is the risk of the quadratic loss $L(\hat{\theta}_k, \theta) = \|\hat{\theta}_k - \theta\|_2^2$ under the Bayes estimator $\hat{\theta}_k(D_k) = \mathbb{E}[\theta \mid D_k]$ — that is, the posterior mean, which is exactly what this code reports as `posterior_mean`.

For the single-parameter, experiment-limited case implemented here, the trace is over a 1×1 covariance and the expectation is estimated by averaging over true ω values. **So the Bayes risk is precisely the mean final posterior variance** that `pipeline/evaluate.py` prints as `final_variance_mean` and that the Baseline results table ranks methods by. They are the same number; the paper’s name for it is the Bayes risk.

Note the framing: the paper is explicit that the *frequentist* route to experiment design goes “with the Cramér–Rao bound formalism by maximizing the quantum Fisher information”, and it deliberately takes the Bayesian route instead. Fisher information is a useful sanity check here (below) but it is **not** the objective.

2.5 The reward

The reward is the **reduction in traced posterior covariance** at each step — Fiderer Eq. (2), implemented at `modules/rewards.py:11-12` and inlined at `train_TRPO_baseline.py:300,317`:

$$R(D_k) = \text{tr Cov}(\theta \mid D_{k-1}) - \text{tr Cov}(\theta \mid D_k) \xrightarrow{d=1} V_{k-1} - V_k$$

with $V_k = \sum_i w_i (\omega_i - \bar{\omega})^2$. This is not an arbitrary choice, and it is **not** a weaker substitute for an information-gain reward. The paper’s reasoning: the negative Bayes risk is the obvious reward but

is far too slow to compute inside a training loop, whereas Eq. (2) is cheap in the SMC framework *and* telescopes to it. Eq. (3), for N experiments per episode:

$$\mathbb{E}_{D_N} \left[\sum_{j=1}^N R(D_j) \right] = \text{const} - r(h \mid p(\theta)), \quad \text{const} = \text{tr Cov}(\theta)$$

So maximising the episode return **provably minimises the Bayes risk** — but the identity holds only for an *undiscounted* return. The paper states this explicitly: “RL has the goal to maximize the expected discounted reward which equals the left-hand side of Eq. (3) **because we set the discount factor to one.**” This repo sets $\gamma = 0.99$, which breaks it — see gotcha 4.

2.6 Aside: why long interrogation times stop helping

Not part of the paper’s framework — included because it explains the shape of the results. Differentiating the likelihood gives the single-shot Fisher information (`modules/policies.py:fisher_information`):

$$I(\omega; t) = \frac{(\partial_\omega p_0)^2}{p_0(1-p_0)} = \frac{t^2 v^2 \sin^2(\omega t)}{1 - v^2 \cos^2(\omega t)}, \quad v = e^{-t/T_2}$$

Without decoherence ($v \rightarrow 1$) this is just $I = t^2$: information grows quadratically in interrogation time, which is why longer is better and why adaptive schemes that grow t pay off. With finite T_2 the envelope $t^2 e^{-2t/T_2}$ peaks at $t = T_2$, so information per shot is **bounded** at $I_{\max} \approx 0.135 T_2^2$. Numerically, for $T_2 = 10$ the optimum sits at $t^* \approx 8.4$ – 9.5 depending on ω (`optimal_t_grid`).

This matches the paper’s observation that “times which exceed T_2 tend to yield no information, which explains why the Bayes risk saturates for the exp-sparse heuristic.”

2.7 The competing constraint: phase ambiguity

Fisher information is a *local* quantity, and maximising it is the wrong move while the prior is still broad. The likelihood depends on ω only through $\cos(\omega t)$, which is periodic. Over the prior support $\Delta\omega = 1$ the fringe wraps once $t > 2\pi \approx 6.28$: several well-separated ω then produce the same phase, the posterior goes **multimodal**, and the filter cannot resolve the aliases.

So a good policy must trade off:

- **small** t — unambiguous, but low information per shot ($I \sim t^2$);
- $t \approx T_2$ — maximum information per shot, but aliases a broad prior;
- $t \gg T_2$ — decohered, zero information, strictly wasted.

The resolution is to grow t *as the posterior narrows*, keeping the phase spread across the posterior at roughly one radian — exactly what the σ^{-1} and PGH heuristics do, and what an RL policy is expected to discover.

Because those heuristics only reach large t *after* the posterior has collapsed, they never alias, which is why the paper discusses large- t failure purely in terms of decoherence. A **constant** $t = 10$, by contrast, aliases from the very first shot — which is why it performs so badly in the measured

results despite sitting at the Fisher optimum. The two statements are consistent; they describe different situations.

3 The inference layer: SMC particle filter

`modules/algorithms/seq_montecarlo.py`

ω is a **static** parameter — it does not evolve — so the filter only ever reweights particles; the particle locations move solely during resampling.

Representation. N particles $\{\omega_i\}$ with log-weights $\{\log w_i\}$. Shapes are $(N, 1)$ / $(N,)$ in the NumPy path, $(B, N, 1)$ / (B, N) in the batched Torch path, where B indexes independent experiments.

Bayes update (`seq_montecarlo.py:97-117`). Log-domain throughout:

$$\log w_i \leftarrow \log w_i + \log P(d \mid \omega_i, t) - \text{logsumexp}(\cdot)$$

Renormalised every step, so weights always sum to one.

Effective sample size (`seq_montecarlo.py:19-21`). The Kish ESS, $\text{ESS} = 1 / \sum_i w_i^2 \in [1, N]$.

Resampling fires when $\text{ESS} < 0.75N$ — an aggressive threshold ($0.5N$ is the usual default), so it triggers on most steps. Two schemes exist; the default everywhere is **Liu–West kernel resampling** (`seq_montecarlo.py:32-71`, $a = 0.98$):

$$\omega'_i = a \omega_{\text{idx}(i)} + (1 - a) \bar{\omega} + \varepsilon_i, \quad \varepsilon_i \sim \mathcal{N}(0, h^2 \Sigma), \quad h^2 = 1 - a^2$$

The shrinkage toward the weighted mean exactly cancels the variance added by the jitter ($a^2 \Sigma + h^2 \Sigma = \Sigma$), so the posterior’s first two moments are preserved. The jitter is what prevents particle impoverishment — plain multinomial resampling (`resample`, `seq_montecarlo.py:24-29`) would collapse all particles onto a handful of duplicated values, since a static parameter never gets diffused by process noise.

4 The control problem

Framed as a Gymnasium environment, `AdaptiveSMCEnv` (`train_TRPO_baseline.py:225-343`):

observation	<code>[posterior_mean, posterior_var, *last_30_interrogation_times] → Box(32,)</code>
action	interrogation time $t \rightarrow \text{Box}(0.1, 3000.0, (1,))$
reward	$V_{k-1} - V_k$
episode	100 measurements (training) / 125 (evaluation), one true ω per episode

5 Methods implemented

5.1 Learned policies

method	optimiser	network	entry point
TRPO	trust-region policy optimisation (<code>sb3_contrib</code>)	<code>MlpPolicy</code> , π and V both [256, 256]	<code>pipeline/train_TRPO_baseline.py</code>
CEM	cross-entropy method over the flat 545-dim weight vector	<code>TimePolicy_Fiderer_16</code> — one hidden layer of 16, <code>Tanh</code> , sigmoid-squashed output	<code>pipeline/train_CEM.py</code>

The CEM network squashes its output structurally, $t = t_{\min} + (t_{\max} - t_{\min}) \sigma(z)$ (`models/nn.py:97`), so it can never emit an out-of-range time. The TRPO actor is an unbounded Gaussian whose sample is clipped in `step()` (`train_TRPO_baseline.py:298`).

5.2 Training-free baselines

`modules/policies.py` — these exist so the learned policies can be measured against something. All share one batched interface and run through **identical** SMC and measurement code via `modules/rollout_generic.py`, which is the only way a comparison means anything.

policy	rule	why it's here
<code>PGHPolicy</code>	$t = 1/ \omega_1 - \omega_2 $, two particles drawn from the posterior	the standard baseline in this literature (Wiebe & Granade); keeps phase spread ≈ 1 rad
<code>PGHPolicy(t_cap=T2)</code>	as above, capped at the coherence time	the standard fix for the decoherence-limited case
<code>SigmaInversePolicy</code>	$t = k/\sigma$	deterministic cousin of PGH, no sampling jitter
<code>FixedTPolicy</code>	constant t	non-adaptive control — separates “adaptivity helps” from “a good t helps”
<code>ExponentialSweepPolicy</code>	$t_k = t_0 r^k$	open-loop ladder; depends on step index, not on data
<code>RandomTPolicy</code>	$t \sim U(0.1, 3000)$	the do-nothing floor

6 Repository layout

<code>modules/</code>	
<code>simulation.py</code>	measurement model (Born rule + T2 decay), NumPy and Torch
<code>rewards.py</code>	posterior variance and variance-reduction reward
<code>algorithms/</code>	
<code>seq_montecarlo.py</code>	SMC filter: init, Bayes update, ESS, Liu-West resampling
<code>CEM.py</code>	cross-entropy method optimiser
<code>rollout.py</code>	episode rollout for a raw torch policy → used by CEM

rollout_sb3.py	episode rollout for an SB3 model, CPU	-> used by
↳ test_TRPO_N		
rollout_sb3_cuda.py	batched rollout for an SB3 model, GPU	-> used by
↳ test_TRPO_N_cuda		
rollout_generic.py	batched rollout for ANY policy	-> used by
↳ evaluate		
policies.py	baseline policy zoo + Fisher-information reference	
models/nn.py	policy networks (TimePolicy_Fiderer_16 is the live one)	
pipeline/		
evaluate.py	evaluate ANY policy, write plots + summary	<- start here
train_TRPO_baseline.py	TRPO training (defines AdaptiveSMCEnv)	
train_CEM.py	CEM training	
test_TRPO_N.py	evaluate a TRPO policy over N true omegas, CPU	
test_TRPO_N_cuda.py	same, batched on GPU	
test_CEM.py	evaluate a CEM policy, single omega	
test_CEM_N.py	evaluate a CEM policy over N true omegas	
	(the four test_*.py are legacy; see below)	
utils/		
plotstyle.py	shared matplotlib style: palette, marks, log axes	
git_utils/git.py	commit / branch / dirty-flag provenance for MLflow	
cov.py, network_fill.py	unused	
artifacts/	scratch output dir, shared across ALL runs (see gotcha 6)	
mlruns/	MLflow artifact root	
mlflow.db	MLflow sqlite backend (~320 MB, gitignored)	

Note there are **four** independent copies of the SMC episode loop: the three `rollout*.py` files plus a fourth inlined in `AdaptiveSMCEnv.step`. They differ in small ways (see gotcha 5).

7 Installation

```
conda create -n NNBQE python=3.11 -y
conda activate NNBQE
pip install -r requirements.txt

# CUDA build of torch, if you have a GPU (skip on Apple Silicon)
pip install torch torchvision --extra-index-url https://download.pytorch.org/whl/cu130
```

Verified working set: Python 3.11.14, torch 2.9.1, stable-baselines3 2.7.1, sb3-contrib 2.7.1, gymnasium 1.2.3, numpy 2.3.5, mlflow 3.8.1. The saved TRPO checkpoint was produced with exactly these versions.

A GPU is **not** required. The batched rollout does ~4.4 ms per (125-step, 2000-particle) episode on CPU, so a full 10 000- ω evaluation takes about a minute.

8 Running things

All commands are run as modules **from the repository root**, with the NNBQE environment active.

Two entry points matter:

<code>python -m pipeline.train_TRPO_baseline / train_CEM</code>	train a policy
<code>python -m pipeline.evaluate -policy <name></code>	evaluate any policy and write plots

`pipeline/evaluate.py` handles **every** method — both learned policies and all six training-free baselines — through identical inference and measurement code, which is what makes the numbers comparable. It supersedes the older `pipeline/test_*.py` scripts (see below).

8.1 Quick check that everything works

No training required — the repo ships a trained TRPO policy, and the baselines need none. This takes a few seconds:

```
python -m pipeline.evaluate --policy trpo \  
    --checkpoint artifacts/trpo_sb3_policy.zip --n-omegas 256  
python -m pipeline.evaluate --policy pgh-capped --n-omegas 256
```

8.2 TRPO, end to end

1. Train.

```
python -m pipeline.train_TRPO_baseline
```

Defaults: 10^5 episodes $\times$ 100 steps = 10^7 timesteps, 10 000 particles, `net_arch=[256,256]`, `target_kl=1e-2`. Expect **4–25 hours** on CPU (the spread is real: the 100k-timestep run took 154 s, the 100M-timestep run averaged $6\times$ slower per step). Set `DEVICE=cuda` if you have a GPU.

Checkpoints are written every 50 000 timesteps, so you can evaluate a partially trained policy at any point and a crash costs you at most one interval:

<code>artifacts/checkpoints/trpo_step_<N>.zip</code>	numbered history
<code>artifacts/trpo_sb3_policy_latest.zip</code>	rolling latest
<code>artifacts/trpo_sb3_policy.zip</code>	final, written when training completes
<code>artifacts/checkpoints/trpo_interrupted.zip</code>	written on crash or Ctrl-C

Shorten a run with the environment overrides:

```
# ~10 minutes instead of hours
N_TRAIN_EPISODES=5000 N_PARTICLES=2000 CHECKPOINT_EVERY_STEPS=25000 \
python -m pipeline.train_TRPO_baseline
```

2. Evaluate.

```
python -m pipeline.evaluate --policy trpo \
    --checkpoint artifacts/trpo_sb3_policy_latest.zip \
    --n-omegas 10000
```

Any of these checkpoints work as `-checkpoint`:

artifacts/trpo_sb3_policy.zip	# best available (run ↪ caaabc8)
artifacts/checkpoints/trpo_step_000250000.zip	# any mid-training ↪ checkpoint
mlruns/3/809c8bf6c4cb45569dc7719cc82e24a9/artifacts/trpo_sb3_policy.zip	

8.3 CEM, end to end

1. Train.

```
python -m pipeline.train_CEM
```

Defaults: 100 generations $\times$ population 1000, elite fraction 0.1, 2000 particles, 100-step episodes, optimising the flat 545-dim weight vector of `TimePolicy_Fiderer_16`. **About 45 minutes** on CPU (~27 s per generation).

Checkpoints every 5 generations:

artifacts/checkpoints/cem_gen_<N>.pt	numbered history
artifacts/cem_policy_latest.pt	rolling latest
artifacts/checkpoints/cem_final.pt	written when training completes

Each `.pt` holds the policy `state_dict` alongside the CEM distribution (`mu`, `sigma`), so a run can be inspected or resumed.

```
# ~2 minutes instead of 45
CEM_GENERATIONS=10 CEM_POP=200 python -m pipeline.train_CEM
```

2. Evaluate.

```
python -m pipeline.evaluate --policy cem \
    --checkpoint artifacts/cem_policy_latest.pt \
    --n-omegas 10000
```

The CEM checkpoint that ships in `mlruns/1/` is from a 2-generation run and is effectively untrained. Train first before drawing any conclusion about CEM.

8.4 Baselines (no training)

```
python -m pipeline.evaluate --policy pgh                # t = 1/|w1 - w2|
python -m pipeline.evaluate --policy pgh-capped         # capped at T2
python -m pipeline.evaluate --policy pgh-capped --t-cap 2.0 # capped elsewhere
python -m pipeline.evaluate --policy sigma-inv --k 1.0   # t = k/sigma
python -m pipeline.evaluate --policy fixed --fixed-t 1.0 # constant t
python -m pipeline.evaluate --policy exp-sweep --t0 0.1 --r 1.05
python -m pipeline.evaluate --policy random             # t ~ U(0.1, 3000)
```

8.5 What evaluation produces

Output goes to `validation/run<N>/` (auto-incremented), or `-out DIR`:

file	contents
<code>summary.json</code>	final variance mean/median/p10/p90, mean & max t , ESS, resample count, runtime
<code>trajectories.npz</code>	raw <code>var</code> , <code>mean</code> , <code>ess</code> , <code>t</code> , <code>omegas</code> , <code>n_resample</code> arrays
<code>posterior_variance_log.png</code>	posterior collapse, log axis — the headline plot
<code>posterior_variance.png</code>	same, linear axis
<code>predicted_t.png</code>	chosen interrogation time vs step, against the T_2 line
<code>ess.png</code>	effective sample size, against the resample threshold
<code>reward_vs_step.png</code>	per-measurement variance reduction
<code>final_variance_vs_omega.png</code>	final variance against true ω — exposes aliasing

Useful flags: `-n-omegas` (default 2000), `-episode-len` (125), `-n-particles` (2000), `-batch-size` (512), `-seed` (42), `-device`, `-out`. Run `python -m pipeline.evaluate -help` for the full list.

Cost is roughly 4.4 ms per episode, so 10 000 ω takes about a minute on CPU.

8.6 Comparing methods

Give each run its own directory and read the summaries side by side:

```

for m in trpo pgh pgh-capped sigma-inv random; do
    extra=""
    [ "$m" = trpo ] && extra="--checkpoint artifacts/trpo_sb3_policy.zip"
    python -m pipeline.evaluate --policy "$m" $extra \
        --n-omegas 4096 --out "validation/cmp_$m"
done

python - <<'EOF'
import json, glob
rows = [json.load(open(f)) for f in sorted(glob.glob("validation/cmp_*/summary.json"))]
print(f"{'method':<32}{'final var':>12}{'median t':>10}")
for r in sorted(rows, key=lambda r: r["final_variance_mean"]):
    print(f"{r['label']:<32}{r['final_variance_mean']:>12.3e}{r['median_t']:>10.2f}")
EOF

```

8.7 Plot styling

All figures use `utils/plotstyle.py` — one validated categorical palette, thin lines, small point marks, and a recessive grid. Import `apply_style()` before creating figures and use `LINE` for trajectories, `log_y()` for log axes (it labels the minor ticks, which matplotlib does not by default), and `refline()` for annotation lines. Change the look in one place rather than per script.

8.8 Loading a policy in your own code

```

from sb3_contrib import TRPO
model = TRPO.load("artifacts/trpo_sb3_policy.zip", device="cpu")

from modules.policies import CEMPolicy, TRPOPolicy, PGHPolicy
cem = CEMPolicy.from_checkpoint("artifacts/cem_policy_latest.pt")

```

Then roll any of them out directly:

```

import numpy as np
from modules.rollout_generic import evaluate_policy
from modules.simulation import FIXED_T2

omegas = np.random.default_rng(42).uniform(0, 1, 4096)
for policy in [TRPOPolicy(model), cem, PGHPolicy(t_cap=FIXED_T2)]:
    r = evaluate_policy(policy, omegas, n_particles=2000, episode_len=125)
    print(f"{policy.label:28s} final var = {r['var'][:, -1].mean():.3e}")

```

`evaluate_policy` returns numpy arrays: `var`, `mean`, `ess` of shape $(B, \text{episode_len} + 1)$ and `t` of shape $(B, \text{episode_len})$. Index k of `var` is the posterior **after** k measurements, so `var[:, 0]` is the prior and `var[:, -1]` is the true final posterior.

8.9 Environment overrides for training

TRPO: `N_PARTICLES`, `EPISODE_LEN`, `HISTORY_SIZE`, `RANDOM_SEED`, `N_TRAIN_EPISODES`, `PLOT_COUNT`, `CHECKPOINT_EVERY_STEPS`, `DEVICE`, `MLFLOW_TRACKING_URI`.

CEM: `N_PARTICLES`, `EPISODE_LEN`, `CEM_POP`, `CEM_ELITE_FRAC`, `CEM_INIT_STD`, `CEM_GENERATIONS`, `HISTORY_SIZE`, `RANDOM_SEED`, `CHECKPOINT_EVERY_GENS`, `MLFLOW_TRACKING_URI`.

```
# 30-second smoke tests
N_PARTICLES=200 EPISODE_LEN=10 N_TRAIN_EPISODES=300 PLOT_COUNT=3 \
CHECKPOINT_EVERY_STEPS=1000 python -m pipeline.train_TRPO_baseline

N_PARTICLES=200 EPISODE_LEN=10 CEM_POP=20 CEM_GENERATIONS=3 \
python -m pipeline.train_CEM
```

8.10 The legacy pipeline/test_*.py scripts

Superseded by `pipeline/evaluate.py`, and **none of them run unmodified**. Each hardcodes an MLflow `RUN_ID` from a machine that no longer exists, and the TRPO ones hardcode `file:///home/chrzanowski/mlflow_tracking` as the tracking store. Additionally:

- `test_CEM.py` calls `rollout()` without the required `resample_fn` argument and raises `TypeError` regardless — it was never updated when the signature changed.
- `test_TRPO_N_cuda.py` hardcodes `DEVICE = "cuda"` and calls `torch.cuda.get_device_name(0)`, so it cannot run without CUDA.
- Both shell wrappers in `scripts/` use `taskset`, which is Linux-only.

To revive one anyway, edit its `RUN_ID` and `mlflow.set_tracking_uri(...)`, or replace the MLflow lookup with a direct path to a local checkpoint.

9 Experiment tracking

Both trainers default to the repo-local sqlite backend and honour `MLFLOW_TRACKING_URI`:

```
mlflow ui --backend-store-uri sqlite:///mlflow.db
```

Experiments: `fiderer / omega_estimation / trpo` and `fiderer / omega_estimation / cem`. (CEM previously logged into the experiment literally named `... / trpo`; that is fixed.)

Logged per episode (TRPO) or per generation (CEM): `mean_reward`, `mean_init_var`, `mean_final_var`, `mean_ess`, `mean_t`, `true_omega`, plus `max_reward`, `sigma_mean`, `mu_norm` for CEM. Params capture the full hyperparameter set plus git commit / branch / dirty flag.

Historical runs already in `mlflow.db` (no retraining needed to replot them):

run	experiment	episodes logged	best mean_final_var	weights on disk
16eda430...	trpo_baseline	270 764	9.28e-4	X crashed before saving
caaabcb8...	trpo_baseline	1 003	2.85e-3	✓ artifacts/trpo_sb3_policy.zip
809c8bf6...	trpo_baseline	100	2.44e-3	✓ in mlruns/3/
b28a7dbc...	trpo_baseline	102	8.19e-2	✓ in mlruns/3/ (barely learned)
bc96386c...	cem	2 generations	—	✓ mlruns/1/models/.../model.pth (untrained)

Extract a metric history without loading MLflow:

```
sqlite3 -readonly mlflow.db \
  "select step, value from metrics
  where run_uuid='16eda43052fd436e8ff76d06ba9d3f0e' and key='mean_final_var'
  order by step;"
```

10 Baseline results

Every method run through identical SMC and measurement code via `pipeline/evaluate.py`. **2048 true ω , 125 measurements, 2000 particles, seed 42, $T_2 = 10$** — the same ω list for every method, so the comparison is paired. Prior variance is $1/12 = 0.0833$.

The **median final var** / **mean** columns are the **Bayes risk** of Eq. (1) — for this single-parameter, experiment-limited problem the traced posterior covariance averaged over true ω is exactly the final posterior variance. Lower is better.

method	median final var	mean	p10	p90	median t	max t
fixed $t = 0.5$	4.09e-04	4.30e-04	3.05e-04	5.80e-04	0.50	0.5
fixed $t = 1$	4.33e-04	4.59e-04	2.98e-04	6.56e-04	1.00	1.0
exponential sweep ($0.1 \cdot 1.05^k$)	1.75e-03	2.04e-03	1.06e-03	3.24e-03	2.06	42.4
TRPO (learned)	2.03e-03	2.79e-03	9.26e-04	5.48e-03	0.75	4.1
PGH (capped at 2)	2.03e-03	2.04e-03	9.41e-04	3.09e-03	2.00	2.0
fixed $t = 2$	2.06e-03	2.07e-03	9.74e-04	3.10e-03	2.00	2.0
PGH (capped at T_2)	3.72e-03	8.39e-03	1.44e-03	1.98e-02	10.00	10.0
PGH (uncapped)	5.19e-03	8.63e-03	2.46e-03	1.61e-02	11.44	3000.0
$1/\sigma$	8.86e-03	1.31e-02	2.99e-03	2.92e-02	8.85	26.3
fixed $t = 10$	8.02e-02	8.08e-02	6.76e-02	9.48e-02	10.00	10.0
random $t \sim U(0.1, 3000)$	8.29e-02	8.07e-02	6.28e-02	8.71e-02	1502	3000.0

Reproduce with:

```
for spec in "fixed --fixed-t 0.5" "fixed --fixed-t 1.0" "pgh" "random"; do
```

```
python -m pipeline.evaluate --policy $spec --n-omegas 2048 --seed 42
done
python -m pipeline.evaluate --policy trpo \
    --checkpoint artifacts/trpo_sb3_policy.zip --n-omegas 2048 --seed 42
```

10.1 What this says

1. The learned policy loses to a constant. TRPO lands mid-table at $2.03\text{e-}03$ while a constant $t = 1$ reaches $4.33\text{e-}04$. Because every method saw the identical ω list, the comparison can be made per- ω : **fixed $t = 1$ beats TRPO on 99.85% of the 2048 frequencies.** This is not sampling noise. TRPO also has by far the widest spread of any adaptive method (p10 $9.26\text{e-}04$ to p90 $5.48\text{e-}03$), which is the signature of an unstable policy rather than a merely suboptimal one — **35.6% of its actions sit exactly on the $t = 0.1$ floor**, in consecutive stretches with a median length of 22 measurements, visible as flat runs in `predicted_t.png`.

Note this is the best *loadable* checkpoint (run `caaabcb8`, 1000 episodes). The 68-hour run reached a $3\times$ better training metric but its weights were lost (gotcha 7), so this comparison cannot speak for a fully trained policy.

2. Large t is catastrophic, and that part is well understood. Constant $t = 10$ ($8.02\text{e-}02$) barely improves on the prior (0.0833) despite sitting exactly at the Fisher-information optimum $t = T_2$. The reason is phase ambiguity: the likelihood depends on ω only through $\cos(\omega t)$, and over the prior support $\Delta\omega = 1$ the fringe wraps once $t > 2\pi \approx 6.28$. Every measurement is then consistent with several well-separated ω , the posterior goes multimodal, and the filter cannot resolve it. **Maximising single-shot Fisher information is exactly the wrong objective while the prior is broad** — which is precisely the tension the physics section describes, and the reason PGH and $1/\sigma$, which both drive t into the aliasing regime, do poorly here.

3. Adaptivity buys nothing in this configuration, which is itself suspicious. The best strategies are the two that never adapt at all. After 125 shots at $t = 1$ the posterior has $\sigma \approx 0.021$, so the aliasing-safe bound has risen to $t < 2\pi/\sigma \approx 300$ — an adaptive schedule *should* be able to exploit that and beat any constant by a wide margin. None does. That points at the setup rather than at adaptivity: the reward goes numerically dead after ~ 20 measurements (gotcha 4), so there is almost no gradient signal for the late-episode behaviour where growing t would pay off.

One caveat on the fine ordering: $t = 0.5$ edging out $t = 1$, and both beating $t = 2$, is *not* explained by aliasing — all three are well inside the unambiguous regime. Liu–West jitter and the aggressive $0.75N$ resample trigger are the likely culprits, but this has not been isolated.

11 Validation over 10 000 ω (the N-scripts)

Both retrained networks, evaluated with the repository’s own multi- ω test harnesses — `pipeline/test_TRPO_N.py` and `pipeline/test_CEM_N.py`. These use the NumPy/float64 rollout (`modules/`

rollout_sb3.py, modules/rollout.py), one ω at a time, and are the **reference implementation**: 10 000 true ω , 125 measurements, 2000 particles, Liu–West resampling.

The ω draw here is independent of the training draw, which is what you want — these are genuinely held-out frequencies.

	TRPO	CEM
training	100 002 episodes, 10^7 timesteps, 2 h 24 m	99 generations $\times$ 1000 pop, 54 min
MLflow run	4719c711 (exp 5)	bf4ad6ae (exp 4)
Bayes risk over 10 000 ω	2.173e-03	2.137e-03
mean t	4.00	6.67
min / max t	0.10 / 9.85	2.16 / 15.60
final ESS (of 2000)	1960	1904


The two methods land in a **statistical tie** (2.17e-03 vs 2.14e-03, a 2% difference) despite completely different optimisers and very different chosen times — CEM settles on interrogation times roughly $1.7\times$ longer than TRPO’s, and never goes below $t = 2.16$, where TRPO ranges down to the action floor.

Reproduce:

```
python -m pipeline.test_TRPO_N      # defaults to run 4719c711
python -m pipeline.test_CEM_N       # defaults to run bf4ad6ae
# override with e.g. RUN_ID=<id> N_OMEGAS=1000 python -m pipeline.test_TRPO_N
```


11.1 Posterior collapse

Mean posterior variance across all 10 000 ω , log scale. Both fall about 1.6 decades from the prior’s 0.0833 and flatten out well before the 125th measurement — the information-starved regime described above.



11.2 Chosen interrogation time


A single representative episode. This is where the two methods differ most.



TRPO oscillates hard between roughly 1 and 7, spiking to ~ 10 in the opening measurements and never settling. CEM is smoother and biased higher. Neither shows the monotone growth in t that the σ^{-1} /PGH analysis says an optimal adaptive policy should produce as the posterior narrows — both are essentially oscillating around a fixed operating point, which is consistent with the constant- t baselines being hard to beat.

Note that both spend time above $t = 2\pi \approx 6.28$, where the fringe aliases over the prior support. That is safe once the posterior is narrow, but TRPO's early spike to $t \approx 10$ happens while the posterior is still wide.

11.3 Effective sample size



ESS stays high (~ 1900 – 1960 of 2000) for both, sawtoothing against the $0.75N = 1500$ resample trigger. The particle filter is healthy; it is not the bottleneck.

11.4 A caveat on the earlier baseline table

The Baseline results table above was produced with `modules/rollout_generic.py`, the batched Torch harness — **not** this NumPy reference path. Cross-checking the two on the same model and the same ω exposed a defect in the batched harness:

harness	mean	median	max	episodes diverging ($>1e-2$)
<code>rollout_sb3</code> (NumPy, float64)	2.18e-03	1.97e-03	7.7e-03	0 of 1000
<code>rollout_generic</code> (Torch, float32)	6.9e-03	2.31e-03	4.8e-01	5–8%

In the failing episodes the particle cloud migrates onto a phase alias outside the prior support (final range e.g. $[0.34, 1.55]$, 99% of particles outside $[0, 1]$) and the posterior variance ends up *larger* than the prior’s — which is impossible for a correct filter.

What has been established: the two filters agree to within 20% when fed identical (t, d) sequences, so the update and resampling code is equivalent; batching is **not** the cause (B=1 with independent seeds diverges at the same rate as B=128); and moving the batched path to float64 halves the divergence rate ($4.7\% \rightarrow 2.0\%$) without eliminating it. The root cause is not yet pinned down.

Consequence: the *median* column of the baseline table is sound — the two harnesses agree there — but the *mean* column is inflated for every method that drives t into the aliasing regime. Treat the medians as the ranking and regard the means as upper bounds pending a re-run through the NumPy path.

12 Divergences from Fiderer et al.

The equations in this repo match the paper. The *setup* around them does not, in several places. These are documented rather than fixed, because changing any of them invalidates the shipped checkpoint and the results table above; they belong in a separate, deliberate change.

	Fiderer	this repo
discount factor γ	1 (required by Eq. 3)	0.99 — gotcha 4
observation	mean, covariance, ≤ 30 past actions, + spent resource	mean, var, 30 past actions — gotcha 3
TRPO initialisation	imitation pretraining, then RL	from scratch (see below)
resource regimes	time-limited and experiment-limited	experiment-limited only (see below)
particles (ω with finite T_2)	2×10^4	10^4 train, 2×10^3 eval
TRPO hidden layers	2×64	[256, 256]
CEM hidden layer	16, ReLU	16, Tanh (<code>models/nn.py:91</code>)
CEM sampling covariance	1/2	<code>CEM_INIT_STD = 1.0</code>
exp-sparse heuristic	$t_k = (9/8)^k$	<code>ExponentialSweepPolicy(t0=0.1, r=1.05)</code>
bound on t	none	<code>T_MAX = 3000</code>

12.1 No imitation pretraining — the leading hypothesis for the gap

Fiderer trains in **two stages**: the network is first initialised by imitation learning on an existing heuristic (σ^{-1} , PGH, or a CEM-trained network), and only then refined with TRPO. The paper is candid that this is a stability measure — “this pretraining step is not strictly necessary but speeds up the training and makes RL more stable” — and every TRPO curve in the paper’s Fig. 2 is pretrained, with the seed heuristic named in the legend.

This repo trains TRPO from scratch. Given that the resulting policy is unstable (35.6% of actions on the action floor) and loses to a constant on 99.85% of frequencies, this is the single most likely explanation for the gap, ahead of the γ bug. It is also the largest piece of missing machinery: a behaviour-cloning stage feeding into TRPO.

12.2 Only the experiment-limited regime is implemented

The paper studies two kinds of episode: **experiment-limited** (a fixed number N of measurements) and **time-limited** (a fixed budget of total evolution time T , so cheap short measurements buy more shots). This repo implements only the former.

That matters for expectations: the paper’s headline claim — an improvement over PGH of “more than one order of magnitude” — is specifically for **time-limited ω estimation with $T_2 = 10$** . **This repo cannot reproduce that result**, because it has no time-limited environment. In the experiment-limited panels the reported NN advantage over PGH is considerably more modest.

12.3 What is *not* the explanation: particle count

The repo evaluates with 2×10^3 particles where the paper uses 2×10^4 for exactly this finite- T_2 case, which looked like a plausible reason for the adaptive heuristics doing poorly — a depleted filter losing the fringe as t grows. **Measured, and it is not.** At 512 ω , seed 42:

heuristic	2 000 particles	20 000 particles
PGH	5.47e-03	5.32e-03
$1/\sigma$	8.34e-03	9.22e-03

A $10\times$ increase moves PGH by 3% and makes $1/\sigma$ marginally worse. Particle depletion is ruled out.

The likelier mechanism is a self-reinforcing aliasing trap. Both heuristics set $t \approx 1/\sigma$; at the variance they plateau near, that is $t \approx 11$, which is past the $t > 2\pi$ aliasing threshold for this prior. The posterior then stays multimodal, σ stays large, t stays ≈ 11 , and the heuristic cannot escape. Capping the time breaks the loop, which is exactly what the results show: PGH capped at 2 reaches 2.04e-03 against 5.19e-03 uncapped. Both heuristics are asymptotically optimal for $T_2 = \infty$, where no such threshold exists.

13 Known issues and gotchas

These are real and worth reading before trusting any number out of this repo.

13.1 1. FIXED_T2 is defined twice

`modules/simulation.py` sets `FIXED_T2 = 100.0` at line 3 and `FIXED_T2 = 10.0` at line 31. **The second wins** — every consumer, including the value logged to MLflow, sees $T_2 = 10$. Line 3 is dead and is a trap for anyone reading the file.

13.2 2. The action range is $300\times$ the coherence time

$t_{\max} = 3000$ with $T_2 = 10$ means $e^{-t/T_2} = e^{-300}$, which underflows to exactly 0.0 in floating point. Any $t \gtrsim 200$ yields $p_0 = 0.5$ identically — a measurement with **literally zero** Fisher information. More than 99% of the action range is physically useless, which makes exploration far harder than it needs to be. `t_max = 3000` would be a sane “coherence-scale bound” only for the dead $T_2 = 100$.

Note this bound is a repo choice, not the paper’s: Fiderer imposes no upper limit on t (its σ^{-1} and PGH heuristics grow without one, and the exp-sparse heuristic grows geometrically forever). A bounded `Box` action space is a reasonable concession to RL, but $O(10\text{--}100)$ would cover everything physically useful. Changing it invalidates comparison with the existing checkpoints.

13.3 3. The observation is missing the resource counter, and is badly scaled

Two separate problems.

Missing input. Fiderer specifies the observation as $\mathbb{E}[\theta \mid D_k]$, $\text{Cov}(\theta \mid D_k)$, the previous actions (at most 30), **and the spent time or number of experiments in the current episode**. This repo builds `[mean, var, *t_history]` (`train_TRPO_baseline.py:287-289`) and omits the resource term entirely. The policy therefore cannot distinguish measurement 5 from measurement 120, which makes any correct endgame behaviour unlearnable — there is no input that tells it the episode is ending. Fixing this means `shape=(3 + history_size,)` and a retrain.

Bad scaling. What is there mixes a variance of order 10^{-4} with interrogation times of order 10^3 , fed unnormalised into the MLP. There is no `VecNormalize` and no `check_env` call.

13.4 4. GAMMA = 0.99 breaks the Bayes-risk identity

`pipeline/train_TRPO_baseline.py:49` sets `GAMMA = 0.99`. The paper sets it to **one**, and that is not incidental — it is what makes the whole training objective correct. Fiderer Eq. (3) holds only for an undiscounted return:

$$\mathbb{E}_{D_N} \left[\sum_{j=1}^N R(D_j) \right] = \text{const} - r(h \mid p(\theta))$$

“RL has the goal to maximize the expected discounted reward which equals the left-hand side of Eq. (3) because we set the discount factor to one.”

With $\gamma = 0.99$ over a 100-step episode, $0.99^{100} \approx 0.37$: late measurements are discounted by nearly two thirds. TRPO is then **not minimising the Bayes risk** — it is minimising a discounted surrogate that overweights early variance reduction. Combined with the fact that the posterior variance collapses by two orders of magnitude within the first ~ 20 measurements (so per-step rewards after that are $\sim 10^{-5}$ regardless), the late-episode signal is close to nonexistent.

Set GAMMA = 1.0 before drawing any conclusion about the method. This is the most likely cause of the greedy, unstable behaviour in the baseline results — 35.6% of actions pinned to the $t = 0.1$ floor.

An earlier version of this README claimed the linear ΔV reward was itself the defect and that the literature uses $-\log V$ instead. That was wrong: the reward matches Fiderer Eq. (2) exactly and is well justified. The discount factor is the actual bug.

13.5 5. The four rollout implementations are not equivalent

- Numerical floors differ: `+1e-12` inside the log (NumPy) vs `clamp(1e-8, 1-1e-8)` on the probability (Torch). CPU and GPU results are not bit-comparable.
- The CUDA path bypasses `model.predict` and reads `mlp_extractor` $\rightarrow$ `action_net` directly, skipping any SB3-side post-processing.
- `mean_t` used to mean different things in different files (the last 30 times vs all of them); `modules/rollout.py` is now fixed to match the env.
- `var_traj` / `var_list` in the legacy paths record the variance **before** each measurement, so their reported `final_posterior_variance` is one measurement stale. `rollout_generic.py` records `episode_len + 1` points and does not have this problem.

13.6 6. artifacts/ is shared scratch, not per-run output

Every run writes plots there and then uploads them to MLflow, so files from different runs interleave by episode number — `posterior_ep_050.png` and `posterior_ep_250000.png` come from different runs. For anything you intend to publish, use the per-run copies under `mlruns/<exp>/<run_id>/artifacts/`.

13.7 7. The best result in the repo is unreproducible

Run `16eda430...` trained for 67.9 hours, reached `mean_final_var = 9.28e-4`, and ended **FAILED** before its single end-of-training `model.save()`. Its metric history and posterior plots survive; **its weights do not**. The best loadable policy (`caaabcb8...`, 1000 episodes) reaches $\sim 2.85e-3$, roughly $3\times$ worse. Closing that gap requires a retrain. Periodic checkpointing has since been added so this cannot recur.

13.8 8. Reproducibility is incomplete

`init_particles` uses the *global* NumPy RNG, so `env.reset(seed=...)` does not control particle initialisation. `TRUE_OMEGAS_LIST` is logged to MLflow as a truncated numpy repr (`'[0.494 0.228 ... 0.520]'`), so the actual ω sequence is recoverable only from `RANDOM_SEED`. Every historical run has `git_dirty = True`, so exact code provenance is not recoverable either.

13.9 9. Dead code

`smc_step` (the pymc path) is called by nothing, yet its `import pymc` at `seq_montecarlo.py:2` drags `pymc/jax/pytensor` into every entry point for zero functionality. `utils/cov.py:posterior_cov_trace` is unused and would raise `ValueError` for the 1-D parameter used here. `models/mn.TimePolicy_Fiderer` has no `forward()` — which means archived checkpoints trained with that class (some older MLflow runs) can no longer be executed.

13.10 10. The shipped TRPO policy is beaten by a one-line constant baseline

Measured, not conjectured — see Baseline results. A constant $t = 1$ reaches a **4.7× lower** median final posterior variance than the trained policy, and wins on 99.85% of individual ω values. Until that gap is closed, no claim that the learned heuristic is doing something useful is supportable.

The candidate causes, in order of suspicion, are all divergences from the paper rather than anything intrinsic to the method: no imitation pretraining, $\gamma = 0.99$ instead of 1 (gotcha 4), and no resource counter in the observation (gotcha 3). The comparison also uses the best *loadable* checkpoint, which trained for 1000 episodes; the 68-hour run scored 3× better before losing its weights (gotcha 7).

13.11 11. `rollout_generic.py` diverges on a few percent of episodes

The batched Torch harness added for `pipeline/evaluate.py` produces catastrophic filter failures at a rate of 5–8%, where the NumPy reference (`modules/rollout_sb3.py`, used by the N-scripts) produces **none in 1000**. In a failing episode the particle cloud locks onto a phase alias outside $[0, 1]$ and the final variance exceeds the prior's.

Established so far: both filters agree when fed identical (t, d) sequences; batching is not the cause; float64 halves the rate but does not remove it. Root cause still open — see the caveat above.

Until this is fixed, prefer `pipeline/test_TRPO_N.py` / `test_CEM_N.py` for any number that matters, and read `pipeline/evaluate.py` medians rather than means.

14 License

See LICENSE.